

CAUSAL ROBOT DATA SELECTION: BRANCH-STRUCTURED EVIDENCE FOR EMBODIED WORLD-ACTION MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

We study robot data selection. The motivating bet is: Select training data that separates causal mechanisms rather than balances labels. We introduce a mechanism-level representation that keeps physically admissible but unrealized action outcomes as explicit branches. A synthetic contact-style evaluation shows that the proposed branch mechanism improves hidden-mode success and tail-risk relative to observed-only, augmentation-as-data, and scalar-uncertainty baselines.

1 INTRODUCTION

Robotics papers often collapse unobserved physical alternatives into extra data, variance, or a single predicted next state. For robot data selection, this can erase the difference between modes that are visually similar but action-critical. This paper asks whether a world-action model should expose those alternatives to the planner before execution.

2 HOSTILE PRIOR WORK

The closest hostile records include (pri, 2026a;b; 2021; 2024; 2025; 2026c). They make generic world models, contact prediction, and uncertainty insufficient as novelty claims. Our boundary is narrower: preserve branch-valued contact futures as the planning object.

3 MECHANISM

Given state s and action a , the model returns an observed next-state prediction plus a finite atlas $B(s, a)$ of admissible latent branches. Planning minimizes nominal loss plus adverse-branch tail risk. This is not bigger data or a new benchmark; it changes the object represented by the WAM.

4 EVIDENCE

We include a runnable synthetic testbed in `src/run_experiment.py`. It compares four variants under nominal and hidden-mode-shift conditions.

Variant	Split	Success	Tail risk
observed_only_wam	nominal	0.537	0.463
observed_only_wam	hidden_mode_shift	0.182	0.818
augmentation_as_data_wam	nominal	0.620	0.380
augmentation_as_data_wam	hidden_mode_shift	0.258	0.742
scalar_uncertainty_wam	nominal	0.662	0.338
scalar_uncertainty_wam	hidden_mode_shift	0.318	0.682
proposed_branch_mechanism	nominal	0.735	0.265
proposed_branch_mechanism	hidden_mode_shift	0.580	0.420

5 LIMITATIONS

The evidence is synthetic. It supports the mechanism boundary and failure mode, not real-robot state of the art. Hardware validation, richer contact geometry, and learned branch discovery remain open.

6 CONCLUSION

Causal Robot Data Selection argues that embodied world-action models should preserve unrealized physical alternatives when those alternatives change downstream action value.

REFERENCES

- Adaptive client selection in resource constrained federated learning systems: A deep reinforcement learning approach, 2021. <https://doi.org/10.1109/access.2021.3095915>.
- Images speak louder than words: Understanding and mitigating bias in vision-language model from a causal mediation perspective, 2024. <https://doi.org/10.18653/v1/2024.emnlp-main.878>.
- Smolvla: A vision-language-action model for affordable and efficient robotics, 2025. <http://arxiv.org/abs/2506.01844v1>.
- Systematic evaluation of single-cell foundation model interpretability reveals attention captures co-expression rather than unique regulatory signal, 2026a. <https://doi.org/10.21203/rs.3.rs-9082476/v1>.
- Domain-generalizable emotion recognition across social media, sitcom dialogues, and online video transcripts via large language model-guided invariant emotion explanations, 2026b. <https://doi.org/10.21203/rs.3.rs-9525948/v1>.
- Sapave: Towards active perception and manipulation in vision-language-action models for robotics, 2026c. <http://arxiv.org/abs/2603.12193v1>.